

(Harvey) Yiyun Fu

CONTACT INFORMATION	<i>Email:</i> harveyfu@uchicago.edu <i>Homepage:</i> harvey-fin.github.io		
RESEARCH INTERESTS	Large Language Models LLM capability evaluation; meta-learning; LLM reasoning efficiency; cross-modality generalization		
EDUCATION	University of Chicago , Chicago, IL	2024 - Present	
	<i>Ph.D. in Computer Science</i> (Advisor: Ari Holtzman)		
	University of Southern California , Los Angeles, CA	2019 - 2023	
	<i>Bachelor of Science in Economics & Mathematics</i>		
RESEARCH EXPERIENCE	University of Chicago , Chicago, IL	Oct 2024 - Present	
	<i>Research Assistant</i> (Advisor: Ari Holtzman)		
	• Large Language Model Reasoning Capability Analysis • Benchmarking Language Model Capabilities Over Long-context Inputs.		
	University of Southern California , Los Angeles, CA	April 2022 - June 2024	
	Allegro Lab & INK Lab <i>Research Assistant</i> (Advisor: Robin Jia and Xiang Ren)		
	• Large Language Model Hallucination Detection I studied LLM hallucination detection problems with a focus on long-sequence generations by inserting contradictory evidence.		
	• Large Language Model Capability Analysis I focused on evaluating the calibration of LLMs through in-context learning and task-level accuracy predictions through a meta-learning process. I also worked on LLM performance predictions through experiment records.		
	University of Southern California , Los Angeles, CA	Spring 2022	
	<i>Student Researcher</i> (Instructor: Goran Muric and Keith Burghardt)		
	• Disinformation Detection I studied falsified information detection for generated texts, images, captions, and also web-scraped data from semantic scholar using a decision tree model.		
BLOG POSTS	Seeing Is Not Reasoning: How VLMs and Their Benchmarks Lean on Text [Link]		
PREPRINTS	Ting-Yun Chang*, Harvey Yiyun Fu* , Deqing Fu, Chenghao Yang, Jesse Thomason, Robin Jia. Value-Aware Stochastic KV Cache Eviction for Reasoning Models. <i>preprint, in submission</i> [Paper]		
	Todd Nief, Harvey Yiyun Fu , Mark Muchane, Ari Holtzman. Subliminal Learning is a LoRA Artifact. <i>preprint, in submission</i> [Paper]		
	Dang Nguyen, Harvey Yiyun Fu , Peter West, Chenhao Tan, Ari Holtzman. Email in the Era of LLMs. <i>preprint, in submission</i> [Paper] [Game]		

PUBLICATIONS	Harvey Yiyun Fu , Aryan Shrivastava, Jared Moore, Peter West, Chenhao Tan, Ari Holtzman. AbsenceBench: Language Models Can't Tell What's Missing. NeurIPS 2025 (Datasets & Benchmarks Track), Spotlight [Paper] [Code]	
	Harvey Yiyun Fu , Qinyuan Ye, Albert Xu, Xiang Ren, Robin Jia. Estimating Large Language Model Capabilities without Labeled Test Data. Findings of EMNLP, 2023 [Paper] [Code]	
	Qinyuan Ye, Harvey Yiyun Fu , Xiang Ren, Robin Jia. How Predictable Are Large Language Model Capabilities? A Case Study on BIG-bench. Findings of EMNLP, 2023 [Paper] [Code]	
INDUSTRY EXPERIENCE	Paragon Partners Asia , Shanghai, China Spring 2022 <i>Financial Analyst</i> Built machine learning models to analyze statistical patterns in cash flow statements.	
	Tempo Asset Management Company , Shanghai, China Summer 2021 <i>Economic Research Intern</i>	
TEACHING EXPERIENCE	CMSC25750: Large Language Models , UChicago Spring 2026 Instructor: Chenhao Tan	
	DATA37712: Foundations of Machine Learning & AI – Part II , UChicago Winter 2026 Instructor: Ari Holtzman	
	CMSC25300: Mathematical Foundations of ML , UChicago Spring 2025, Fall 2025 Instructor: Haifeng Xu and Tian Li	
	ITP115: Programming in Python , USC Aug 2022 - Jan 2024 Instructor: Trina Gregory and Barrett Koster	
AWARDS AND HONORS	USC Provost's Undergrad Research Fellowship Apr 2023 and Sep 2023 USC Academic Achievement Award Jul 2021	

Last update: Jun 15, 2026